

Python Project

Nischal Pant

2023-08-25

The libraries you will use are already loaded for you below

```
In [44]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from itertools import chain
```

First glance

Importing packages

```
In [45]: level_salary = pd.read_csv('Levels_Fyi_Salary_Data.csv')
cost_of_living = pd.read_csv('cost_of_living.csv')
country_codes = pd.read_excel('country_codes.xlsx')
ds_salaries = pd.read_csv('ds_salaries.csv')

print(level_salary.columns)
print(cost_of_living.columns)
print(country_codes.columns)
print(ds_salaries.columns)
```

Index(['timestamp', 'company', 'level', 'title', 'totalyearlycompensation',
 'location', 'yearsofexperience', 'yearsatcompany', 'tag', 'basesalary',
 'stockgrantvalue', 'bonus', 'gender', 'otherdetails', 'cityid', 'dmaid',
 'rowNumber', 'Masters_Degree', 'Bachelors_Degree', 'Doctorate_Degree',
 'Highschool', 'Some_College', 'Race_Asian', 'Race_White',
 'Race_Two_Or_More', 'Race_Black', 'Race_Hispanic', 'Race', 'Education'],
 dtype='object')
Index(['Rank', 'City', 'Cost of Living Index', 'Rent Index',
 'Cost of Living Plus Rent Index', 'Groceries Index',
 'Restaurant Price Index', 'Local Purchasing Power Index'],
 dtype='object')
Index(['Country', 'Alpha-2 code', 'Alpha-3 code', 'Numeric'], dtype='object')
Index(['Unnamed: 0', 'work_year', 'experience_level', 'employment_type',
 'job_title', 'salary', 'salary_currency', 'salary_in_usd',
 'employee_residence', 'remote_ratio', 'company_location',
 'company_size'],
 dtype='object')

```
In [46]: #here we can Look into what the files Look Like and possibly think of ways we can comb
```

```
In [47]: #first lets see what our original dataset looks like
print(cost_of_living.head(10))
```

```
print(cost_of_living.dtypes)

# there is a lot of data so its best to group them together, and we should group them by city
city_country = cost_of_living['City'].str.split(',', expand=True)
cost_of_living['City'] = city_country[0]
cost_of_living['Country'] = city_country[1]
print(cost_of_living.head(10))

# here we can see some columns that we dont really need for our analysis are they are irrelevant
cost_of_living = cost_of_living.drop(columns=['City', 'Rank'])
column_order = ['Country', 'Cost of Living Index', 'Rent Index', 'Cost of Living Plus',
                'Groceries Index', 'Restaurant Price Index', 'Local Purchasing Power']
cost_of_living = cost_of_living[column_order]

# we can now put them in order of whatever we like. Having the city names in the front
```

Rank	City	Cost of Living Index	Rent Index	\
0	NaN Hamilton, Bermuda	149.02	96.10	
1	NaN Zurich, Switzerland	131.24	69.26	
2	NaN Basel, Switzerland	130.93	49.38	
3	NaN Zug, Switzerland	128.13	72.12	
4	NaN Lugano, Switzerland	123.99	44.99	
5	NaN Lausanne, Switzerland	122.03	59.55	
6	NaN Beirut, Lebanon	120.47	27.76	
7	NaN Bern, Switzerland	118.16	46.12	
8	NaN Geneva, Switzerland	114.05	75.05	
9	NaN Stavanger, Norway	104.61	35.38	
Cost of Living Plus Rent Index	Groceries Index	Restaurant Price Index	\	
0	124.22	157.89	155.22	
1	102.19	136.14	132.52	
2	92.70	137.07	130.95	
3	101.87	132.61	130.93	
4	86.96	129.17	119.80	
5	92.74	122.56	127.01	
6	77.01	141.33	116.95	
7	84.39	118.37	120.88	
8	95.77	112.70	126.31	
9	72.16	102.46	107.51	
Local Purchasing Power Index				
0	79.43			
1	129.79			
2	111.53			
3	143.40			
4	111.96			
5	127.01			
6	15.40			
7	112.46			
8	120.60			
9	85.90			
Rank		float64		
City		object		
Cost of Living Index		float64		
Rent Index		float64		
Cost of Living Plus Rent Index		float64		
Groceries Index		float64		
Restaurant Price Index		float64		
Local Purchasing Power Index		float64		
dtype: object				
Rank	City	Cost of Living Index	Rent Index	\
0	NaN Hamilton	149.02	96.10	
1	NaN Zurich	131.24	69.26	
2	NaN Basel	130.93	49.38	
3	NaN Zug	128.13	72.12	
4	NaN Lugano	123.99	44.99	
5	NaN Lausanne	122.03	59.55	
6	NaN Beirut	120.47	27.76	
7	NaN Bern	118.16	46.12	
8	NaN Geneva	114.05	75.05	
9	NaN Stavanger	104.61	35.38	
Cost of Living Plus Rent Index	Groceries Index	Restaurant Price Index	\	
0	124.22	157.89	155.22	
1	102.19	136.14	132.52	
2	92.70	137.07	130.95	

3	101.87	132.61	130.93
4	86.96	129.17	119.80
5	92.74	122.56	127.01
6	77.01	141.33	116.95
7	84.39	118.37	120.88
8	95.77	112.70	126.31
9	72.16	102.46	107.51

	Local Purchasing Power Index	Country
0	79.43	Bermuda
1	129.79	Switzerland
2	111.53	Switzerland
3	143.40	Switzerland
4	111.96	Switzerland
5	127.01	Switzerland
6	15.40	Lebanon
7	112.46	Switzerland
8	120.60	Switzerland
9	85.90	Norway

```
In [48]: level_salary.dropna(inplace=True)
cost_of_living.dropna(inplace=True)
country_codes.dropna(inplace=True)
ds_salaries.dropna(inplace=True)

print(cost_of_living.head(5))

#now that we have grouped together by country names and avraged our index columns. A r
#more organized.

avg_cost_of_living = cost_of_living.groupby('Country').mean().reset_index()
print(avg_cost_of_living.head())

#it seems that some countries are in code so we can now use the country code data and
```

	Country	Cost of Living Index	Rent Index	\
0	Bermuda	149.02	96.10	
1	Switzerland	131.24	69.26	
2	Switzerland	130.93	49.38	
3	Switzerland	128.13	72.12	
4	Switzerland	123.99	44.99	

	Cost of Living Plus Rent Index	Groceries Index	Restaurant Price Index	\
0	124.22	157.89	155.22	
1	102.19	136.14	132.52	
2	92.70	137.07	130.95	
3	101.87	132.61	130.93	
4	86.96	129.17	119.80	

	Local Purchasing Power Index			
0		79.43		
1		129.79		
2		111.53		
3		143.40		
4		111.96		

	Country	Cost of Living Index	Rent Index	\
0	AK	91.23	39.290000	
1	AL	78.82	28.190000	
2	AR	59.26	25.600000	
3	AZ	65.79	34.833333	
4	Afghanistan	21.35	3.170000	

	Cost of Living Plus Rent Index	Groceries Index	Restaurant Price Index	\
0	66.880000	97.950000	78.76	
1	55.090000	84.300000	75.48	
2	43.480000	57.280000	64.63	
3	51.276667	63.963333	69.79	
4	12.830000	15.220000	14.85	

	Local Purchasing Power Index			
0		118.630000		
1		84.930000		
2		131.070000		
3		107.683333		
4		22.790000		

```
In [6]: print(country_codes.columns)
print(country_codes.head(10))
```

#there are 4 variables and now we can find a way to merge country codes and our cost of living index

```
Index(['Country', 'Alpha-2 code', 'Alpha-3 code', 'Numeric'], dtype='object')
      Country Alpha-2 code Alpha-3 code Numeric
0      Afghanistan      AF      AFG        4
1       Albania         AL      ALB        8
2       Algeria         DZ      DZA       12
3  American Samoa      AS      ASM       16
4      Andorra         AD      AND       20
5       Angola          AO      AGO       24
6     Anguilla          AI      AIA      660
7   Antarctica         AQ      ATA       10
8  Antigua and Barbuda      AG      ATG       28
9      Argentina         AR      ARG       32
```

```
In [19]: #the easiest way we can combine the cost of living and country code is by creating a mapping from country codes to country names in our dataset
code_to_name_mapping = dict(zip(country_codes['Alpha-2 code'], country_codes['Country']))

#now we can replace the country codes in avg_cost_of_living with the corresponding country names
avg_cost_of_living['Country'] = avg_cost_of_living['Country'].replace(code_to_name_mapping)

merged_data = pd.merge(avg_cost_of_living, country_codes, on='Country', how='left')
column_order = ['Country', 'Cost of Living Index', 'Cost of Living Plus Rent Index', 'Rent Index', 'Groceries Index', 'Restaurant Price Index', 'Local Purchasing Power Index']
merged_data_cofl = merged_data[column_order]
print(merged_data_cofl.head())

#in the code above, I joined the two datasets together and now we have the indexes and
#there are still two letter country names which suggests that they are actually US states
```

	Country	Cost of Living Index	Cost of Living Plus Rent Index	Rent Index	Groceries Index	Restaurant Price Index	Local Purchasing Power Index
0	AK	91.23	66.880000	39.290000	97.950000	78.76	118.630000
1	Albania	78.82	55.090000	28.190000	84.300000	75.48	84.930000
2	Argentina	59.26	43.480000	25.600000	57.280000	64.63	131.070000
3	Azerbaijan	65.79	51.276667	34.833333	63.963333	69.79	107.683333
4	Afghanistan	21.35	12.830000	3.170000	15.220000	14.85	22.790000

```
In [8]: merged_data_cofl = merged_data_cofl[merged_data_cofl['Country'].str.len() != 2]
print(merged_data_cofl)
```

#here I have just removed the US states that were from the other dataset. Now it looks like this:

	Country	Cost of Living Index	Cost of Living Plus Rent Index	\
1	Albania	78.82	55.090000	
2	Argentina	59.26	43.480000	
3	Azerbaijan	65.79	51.276667	
4	Afghanistan	21.35	12.830000	
5	Albania	38.68	25.860000	
..	
162	Holy See (the)	67.75	53.580000	
163	Venezuela	45.31	29.730000	
164	Vietnam	37.93	27.195000	
167	Zambia	33.57	22.600000	
168	Zimbabwe	45.69	28.750000	
	Rent Index	Groceries Index	Restaurant Price Index	\
1	28.190000	84.300000	75.48	
2	25.600000	57.280000	64.63	
3	34.833333	63.963333	69.79	
4	3.170000	15.220000	14.85	
5	11.330000	30.990000	29.86	
..	
162	37.520000	65.070000	77.50	
163	12.080000	37.600000	48.60	
164	15.030000	39.285000	20.39	
167	10.180000	32.850000	23.63	
168	9.560000	37.050000	39.05	
	Local Purchasing Power Index			
1		84.930000		
2		131.070000		
3		107.683333		
4		22.790000		
5		31.150000		
..		...		
162		130.940000		
163		15.870000		
164		30.770000		
167		37.480000		
168		17.590000		

[150 rows x 7 columns]

```
In [9]: print(ds_salaries.head(5))
print(ds_salaries.columns)
```

#this is the next dataset that we will be filtering and merging with our previously merged visual before proceeding. This is the same dataset which we analysed in our R project

```

      Unnamed: 0  work_year experience_level employment_type \
0          0        2020             MI            FT
1          1        2020             SE            FT
2          2        2020             SE            FT
3          3        2020             MI            FT
4          4        2020             SE            FT

      job_title  salary salary_currency salary_in_usd \
0   Data Scientist    70000           EUR       79833
1 Machine Learning Scientist  260000           USD      260000
2     Big Data Engineer   85000           GBP      109024
3  Product Data Analyst   20000           USD      20000
4  Machine Learning Engineer  150000           USD      150000

employee_residence  remote_ratio company_location company_size
0                  DE          0             DE            L
1                  JP          0             JP            S
2                  GB          50            GB            M
3                  HN          0             HN            S
4                  US          50            US            L
Index(['Unnamed: 0', 'work_year', 'experience_level', 'employment_type',
       'job_title', 'salary', 'salary_currency', 'salary_in_usd',
       'employee_residence', 'remote_ratio', 'company_location',
       'company_size'],
      dtype='object')

```

```

In [49]: #first lets only use the columns that apply to us.
# we want data scientist positions, salary_in_us for comparision, the location of the

ds_salaries_filtered = ds_salaries[['job_title', 'salary_in_usd', 'company_location',
ds_salaries_filtered = ds_salaries_filtered[ds_salaries_filtered['job_title'] == 'Data'
ds_salaries_filtered.reset_index(drop=True, inplace=True)
print(ds_salaries_filtered.head(10))

#this is our first filtered salaries.

#lets create the box plot and first we can define the IQR values and then plug them in
Q1 = ds_salaries_filtered['salary_in_usd'].quantile(0.25)
Q3 = ds_salaries_filtered['salary_in_usd'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

#now lets see what our global salary looks like
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.boxplot(ds_salaries_filtered['salary_in_usd'], vert=False)
plt.title('Boxplot of Salary (USD)')
plt.xlabel('Salary (USD)')

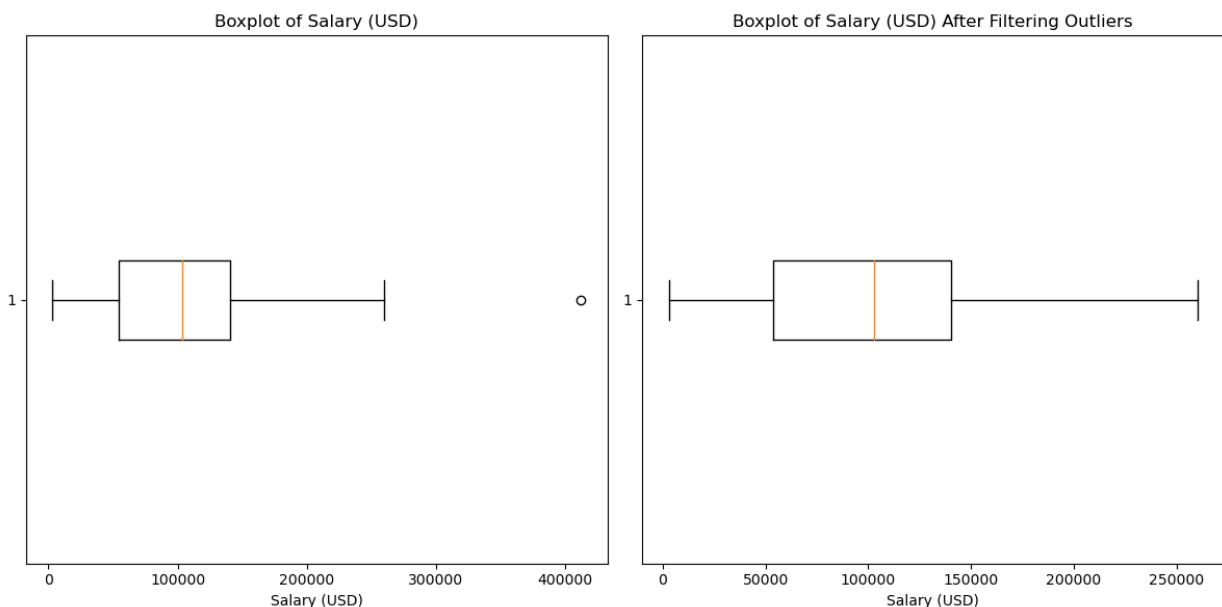
#here we can filter to not show anything past the lower and upper boundries
filtered_salaries = ds_salaries_filtered[
    (ds_salaries_filtered['salary_in_usd'] >= lower_bound) &
    (ds_salaries_filtered['salary_in_usd'] <= upper_bound)
]
plt.subplot(1, 2, 2)
plt.boxplot(filtered_salaries['salary_in_usd'], vert=False)

```

```
plt.title('Boxplot of Salary (USD) After Filtering Outliers')
plt.xlabel('Salary (USD)')
plt.tight_layout()
plt.show()
```

*#it looks like it does make a difference. This was what I was lacking in my R project
#this step in this project to stop outliers from skewing the values too much. And also
#have a great discrepancy.*

	job_title	salary_in_usd	company_location	experience_level
0	Data Scientist	79833	DE	MI
1	Data Scientist	35735	HU	MI
2	Data Scientist	51321	FR	EN
3	Data Scientist	40481	IN	MI
4	Data Scientist	39916	FR	EN
5	Data Scientist	68428	US	SE
6	Data Scientist	45760	US	MI
7	Data Scientist	76958	GB	MI
8	Data Scientist	105000	US	MI
9	Data Scientist	38776	ES	MI



In [50]:

```
#there seems to me multile entries of the same country because they have a slight diff
#states entries
# we can group them together and avrage them out

merged_salaries = ds_salaries_filtered[['company_location', 'salary_in_usd']]
merged_salaries.rename(columns={'company_location': 'Country', 'salary_in_usd': 'Salary (USD)'}, inplace=True)
print(merged_salaries.head(10))

merged_salaries = merged_salaries.groupby('Country')['Salary (USD)'].mean().reset_index()
print(merged_salaries)

#Now lets create a dictionary to replace the two letter code with the alpha-2 code to t
country_code_dict = dict(zip(country_codes['Alpha-2 code'], country_codes['Country']))
merged_salaries['Country'] = merged_salaries['Country'].map(country_code_dict)
print(merged_salaries.head(10))
```

```

Country    Salary (USD)
0        DE        79833
1        HU        35735
2        FR        51321
3        IN        40481
4        FR        39916
5        US        68428
6        US        45760
7        GB        76958
8        US       105000
9        ES        38776
Country    Salary (USD)
0    Austria  76352.000000
1  Australia  86703.000000
2      Brazil 12901.000000
3     Canada  77787.000000
4  Switzerland 122346.000000
5      Chile  40038.000000
6    Germany  69640.142857
7    Algeria 100000.000000
8      Spain  41136.666667
9     France  50085.571429

```

C:\Users\Nischal\AppData\Local\Temp\ipykernel_15080\3062950383.py:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
merged_salaries.rename(columns={'company_location': 'Country', 'salary_in_usd': 'Salary (USD)'}, inplace=True)
```

In []:

In [51]:

```
print(merged_data_cofl.columns)
print(merged_salaries.columns)
```

```
# now we have to group the country with multiple entries in our cofl file aswell as the
# I believe I did this before but my code seems to throw an error so i am putting this
grouped_cofl = merged_data_cofl.groupby('Country').mean()
grouped_salaries = merged_salaries.groupby('Country').mean()
```

```
#now lets merge them into a new one using the country name, so it will have the avg salary
grouped_data = pd.merge(grouped_salaries, grouped_cofl, on='Country', how='left')
grouped_data.dropna(inplace=True)
```

```
(grouped_data.head(10))
```

```
Index(['Country', 'Cost of Living Index', 'Cost of Living Plus Rent Index',
       'Rent Index', 'Groceries Index', 'Restaurant Price Index',
       'Local Purchasing Power Index'],
      dtype='object')
Index(['Country', 'Salary (USD)'], dtype='object')
```

Out[51]:

	Salary (USD)	Cost of Living Index	Cost of Living Plus Rent Index	Rent Index	Groceries Index	Restaurant Price Index	Local Purchasing Power Index
Country							
Algeria	100000.000000	29.840000	18.980000	6.670000	30.250000	20.790000	21.780000
Australia	86703.000000	77.601000	59.633000	39.267000	77.064000	75.104000	106.108000
Austria	76352.000000	72.870000	52.396000	29.194000	66.658000	70.666000	76.552000
Brazil	12901.000000	34.578571	23.137143	10.164286	29.082857	27.838571	29.360000
Canada	77787.000000	76.750160	66.927788	55.794968	77.501346	74.865417	108.086923
Chile	40038.000000	45.640000	31.420000	15.300000	40.120000	47.920000	32.080000
France	50085.571429	77.741667	55.858333	31.056667	77.201667	75.265000	82.336667
Germany	69640.142857	67.281154	49.305385	28.933846	52.963077	65.510769	100.196538
Hungary	35735.000000	41.767500	27.955000	12.300000	37.340000	34.515000	49.812500
India	26108.250000	46.258152	34.393696	20.947065	47.470000	43.115326	85.529783

In [52]: `print(level_salary.head(10))`

```

level_salary = level_salary[(level_salary['title'] == 'Data Scientist') & (level_salary['yearsofexperience'] > 0)]
grouped_stats = level_salary.groupby('yearsofexperience')['totalyearlycompensation'].agg(['mean', 'std', 'min', 'max'])
grouped_stats.rename(columns={'<lambda>': 'IQR'}, inplace=True)
print(grouped_stats)

plt.figure(figsize=(10, 6))
sns.boxplot(x=level_salary['yearsofexperience'], y=level_salary['totalyearlycompensation'])
plt.xlabel('Years of Experience')
plt.ylabel('Total Yearly Compensation')
plt.title('Box Plot of Total Yearly Compensation by Years of Experience')
plt.show()
#I know this is a fyi file but I wanted to see the file in detail and include it in the notebook
#its promising so I will keep putting in the work

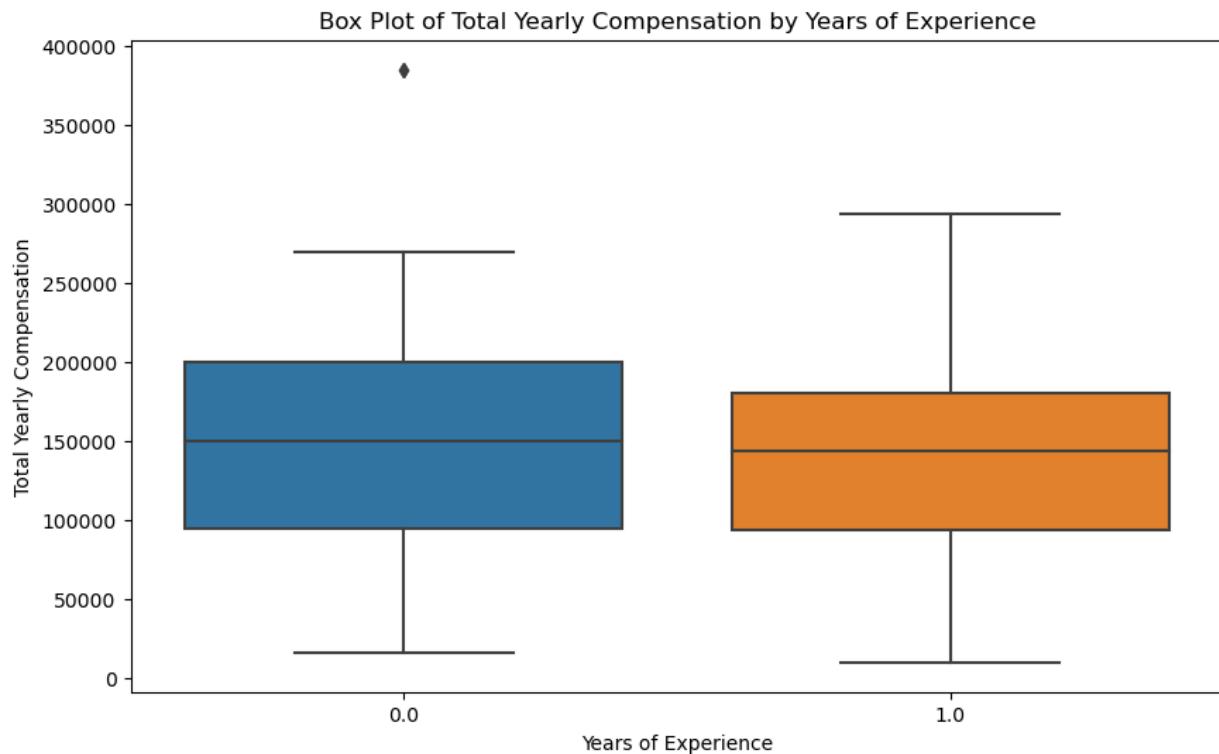
```

	timestamp	company	level	title	\	
15710	1/27/2020 22:59:06	Google	L6	Software Engineer		
23532	7/3/2020 19:56:38	Microsoft	61	Software Engineer		
23533	7/3/2020 20:03:57	Google	L5	Software Engineer		
23534	7/3/2020 20:05:37	Microsoft	62	Software Engineer		
23535	7/3/2020 20:19:06	Blend	IC3	Software Engineer		
23537	7/3/2020 20:24:20	Amazon	L6	Software Engineer		
23538	7/3/2020 20:31:33	Chevron	PSG 20	Software Engineer		
23540	7/3/2020 22:35:58	Amazon	L7	Software Engineering Manager		
23541	7/3/2020 22:51:23	Shopify	L6	Software Engineer		
23543	7/3/2020 23:39:06	Apple	ICT3	Software Engineer		
	totalyearlycompensation		location	yearsofexperience	\	
15710	400000	Sunnyvale, CA		5.0		
23532	136000	Redmond, WA		3.0		
23533	337000	San Bruno, CA		6.0		
23534	222000	Seattle, WA		4.0		
23535	187000	San Francisco, CA		5.0		
23537	310000	Seattle, WA		15.0		
23538	113000	Houston, TX		3.0		
23540	620000	Seattle, WA		19.0		
23541	98000	Toronto, ON, Canada		9.0		
23543	180000	Vancouver, BC, Canada		1.0		
	yearsatcompany	tag	basesalary	...	\	
15710	5.0	Distributed Systems (Back-End)	210000.0	...		
23532	2.0	DevOps	124000.0	...		
23533	6.0	Full Stack	177000.0	...		
23534	4.0	API Development (Back-End)	164000.0	...		
23535	0.0	Full Stack	165000.0	...		
23537	3.0	ML / AI	160000.0	...		
23538	3.0	DevOps	103000.0	...		
23540	7.0	Full Stack	160000.0	...		
23541	4.0	Web Development (Front-End)	78000.0	...		
23543	1.0	ML / AI	130000.0	...		
	Doctorate_Degree	Highschool	Some_College	Race_Asian	Race_White	\
15710	1	0	0	1	0	
23532	0	0	0	0	0	
23533	0	0	0	1	0	
23534	0	0	0	1	0	
23535	0	0	0	0	1	
23537	0	0	0	1	0	
23538	0	0	0	0	0	
23540	0	0	0	1	0	
23541	0	0	0	1	0	
23543	0	0	0	1	0	
	Race_Two_Or_More	Race_Black	Race_Hispanic	Race	\	
15710	0	0	0	Asian		
23532	1	0	0	Two Or More		
23533	0	0	0	Asian		
23534	0	0	0	Asian		
23535	0	0	0	White		
23537	0	0	0	Asian		
23538	0	0	1	Hispanic		
23540	0	0	0	Asian		
23541	0	0	0	Asian		
23543	0	0	0	Asian		

```
Education
15710          PhD
23532 Bachelor's Degree
23533 Bachelor's Degree
23534 Master's Degree
23535 Bachelor's Degree
23537 Bachelor's Degree
23538 Bachelor's Degree
23540 Bachelor's Degree
23541 Bachelor's Degree
23543 Bachelor's Degree
```

[10 rows x 29 columns]

	mean	median
yearsofexperience		
0.0	151000.00000	149500.0
1.0	140219.178082	144000.0



```
In [53]: print(grouped_data.head())
print(grouped_data.columns)
print(grouped_data.dtypes)
```

```

Salary (USD)  Cost of Living Index  Cost of Living Plus Rent Index \
Country
Algeria      1000000.0           29.840000          18.980000
Australia    86703.0            77.601000          59.633000
Austria      76352.0            72.870000          52.396000
Brazil        12901.0            34.578571          23.137143
Canada       77787.0            76.750160          66.927788

Rent Index  Groceries Index  Restaurant Price Index \
Country
Algeria      6.670000           30.250000          20.790000
Australia    39.267000          77.064000          75.104000
Austria      29.194000          66.658000          70.666000
Brazil        10.164286          29.082857          27.838571
Canada       55.794968          77.501346          74.865417

Local Purchasing Power Index
Country
Algeria      21.780000
Australia    106.108000
Austria      76.552000
Brazil        29.360000
Canada       108.086923
Index(['Salary (USD)', 'Cost of Living Index',
       'Cost of Living Plus Rent Index', 'Rent Index', 'Groceries Index',
       'Restaurant Price Index', 'Local Purchasing Power Index'],
      dtype='object')
Salary (USD)          float64
Cost of Living Index  float64
Cost of Living Plus Rent Index  float64
Rent Index            float64
Groceries Index      float64
Restaurant Price Index float64
Local Purchasing Power Index  float64
dtype: object

```

```

In [54]: grouped_data.reset_index(inplace=True)
print(grouped_data.head())

grouped_data.nlargest(5, 'Salary (USD)')

#this is our final grouped data. I've tried to incorporate everything I've learned in
#missing a few things. I am somewhat satisfied with how it turned out so far but could
#the top countries with highest salaries.

```

```

Country    Salary (USD)    Cost of Living Index \
0   Algeria      1000000.0          29.840000
1   Australia     86703.0           77.601000
2   Austria       76352.0           72.870000
3   Brazil        12901.0           34.578571
4   Canada        77787.0           76.750160

Cost of Living Plus Rent Index    Rent Index    Groceries Index \
0                           18.980000    6.670000      30.250000
1                           59.633000   39.267000      77.064000
2                           52.396000   29.194000      66.658000
3                           23.137143   10.164286      29.082857
4                           66.927788   55.794968      77.501346

Restaurant Price Index    Local Purchasing Power Index
0                     20.790000          21.780000
1                     75.104000          106.108000
2                     70.666000          76.552000
3                     27.838571          29.360000
4                     74.865417          108.086923

```

Out[54]:

	Country	Salary (USD)	Cost of Living Index	Cost of Living Plus Rent Index	Rent Index	Groceries Index	Restaurant Price Index	Local Purchasing Power Index
18	Switzerland	122346.0	124.075714	93.802857	59.495714	126.945714	126.914286	122.392857
10	Israel	119059.0	83.118000	65.715000	45.995000	75.532000	87.640000	105.665000
0	Algeria	100000.0	29.840000	18.980000	6.670000	30.250000	20.790000	21.780000
1	Australia	86703.0	77.601000	59.633000	39.267000	77.064000	75.104000	106.108000
4	Canada	77787.0	76.750160	66.927788	55.794968	77.501346	74.865417	108.086923

In [55]: #now lets define a function and then implement our parameters

```

def create_bar_plot(data, x_column, y_column, color_palette):
    plt.figure(figsize=(10, 6))
    sns.barplot(
        data=data,
        x=x_column,
        y=y_column,
        palette=color_palette,
        errorbar=("pi", 50),
        capsize=.4,
        errcolor=".5",
        linewidth=3,
        edgecolor=".5",
        facecolor=(0, 0, 0, 0)
    )
    plt.show()

#now lets create our indexes which we will use to build our boxplots
index_columns = ['Cost of Living Index', 'Rent Index', 'Cost of Living Plus Rent Index']
color_palette = ["red", "green", "blue"]

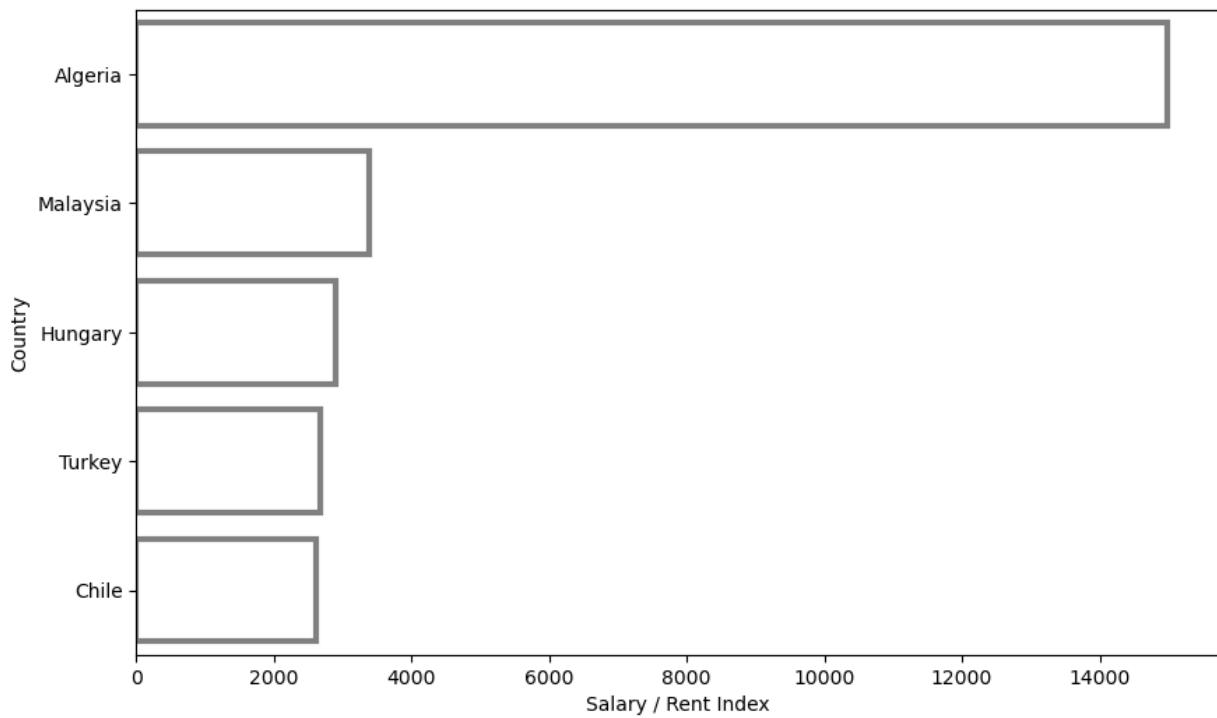
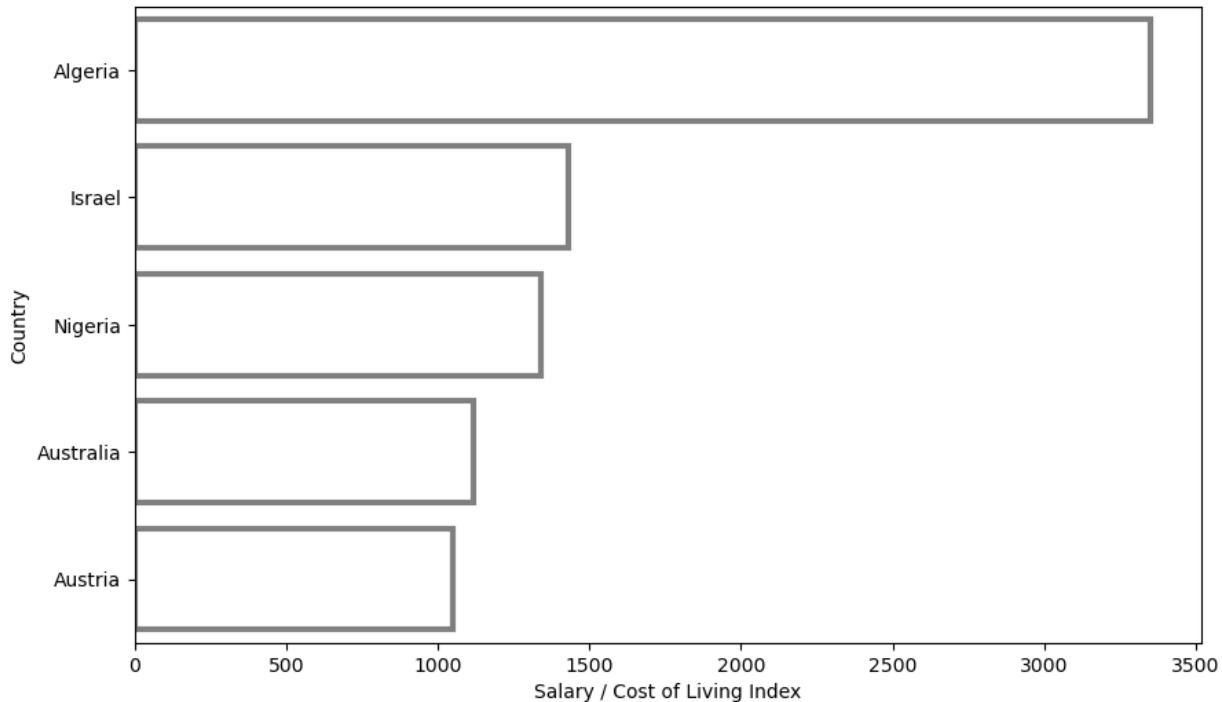
#I dont know why the color palette is not showing. I look online and none of the solutions
#it here until I figure it out.

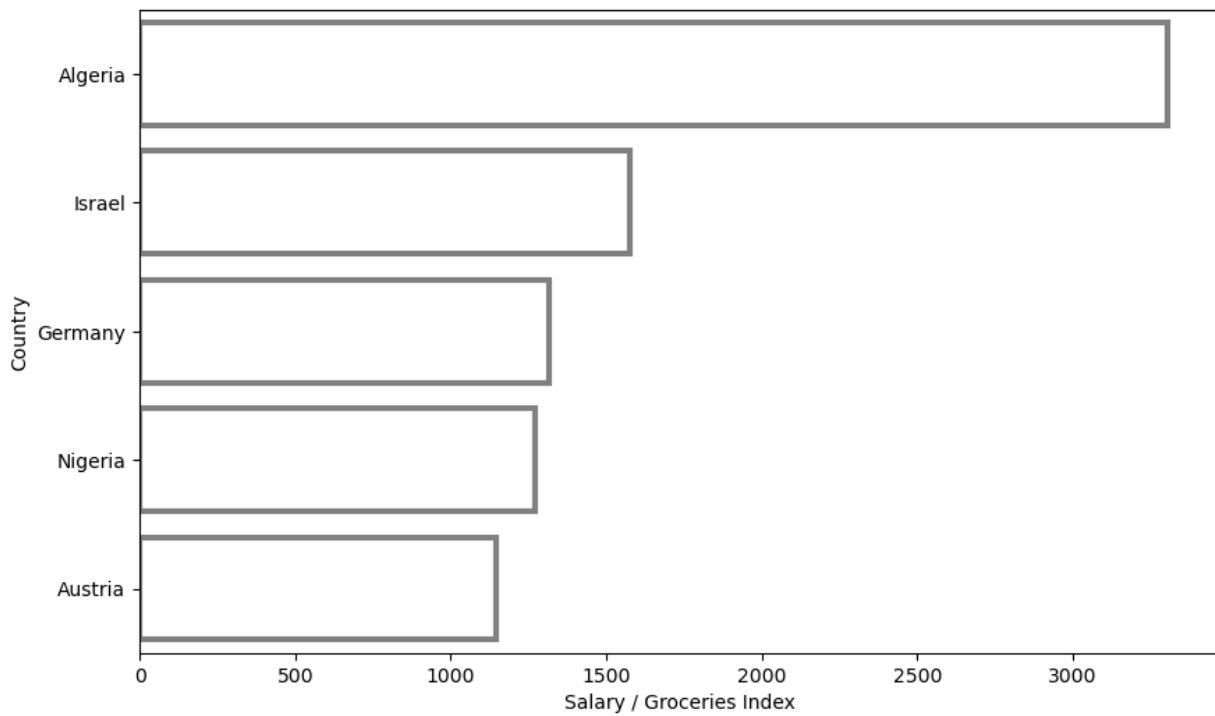
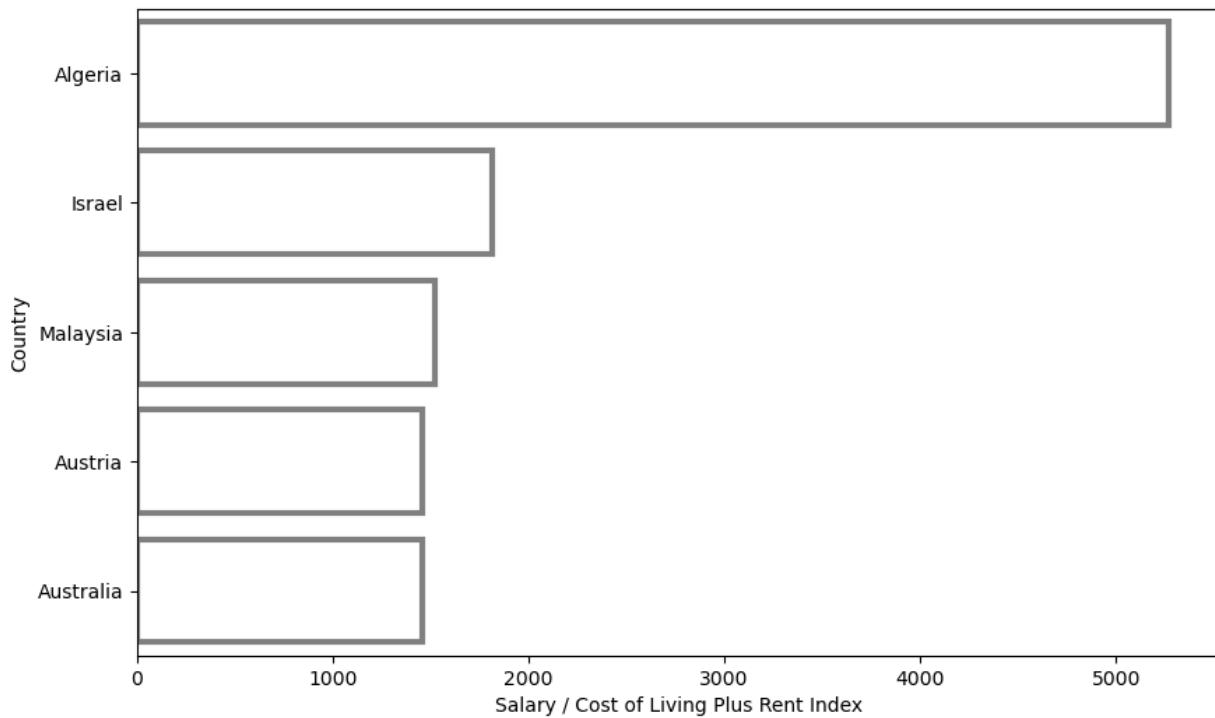
```

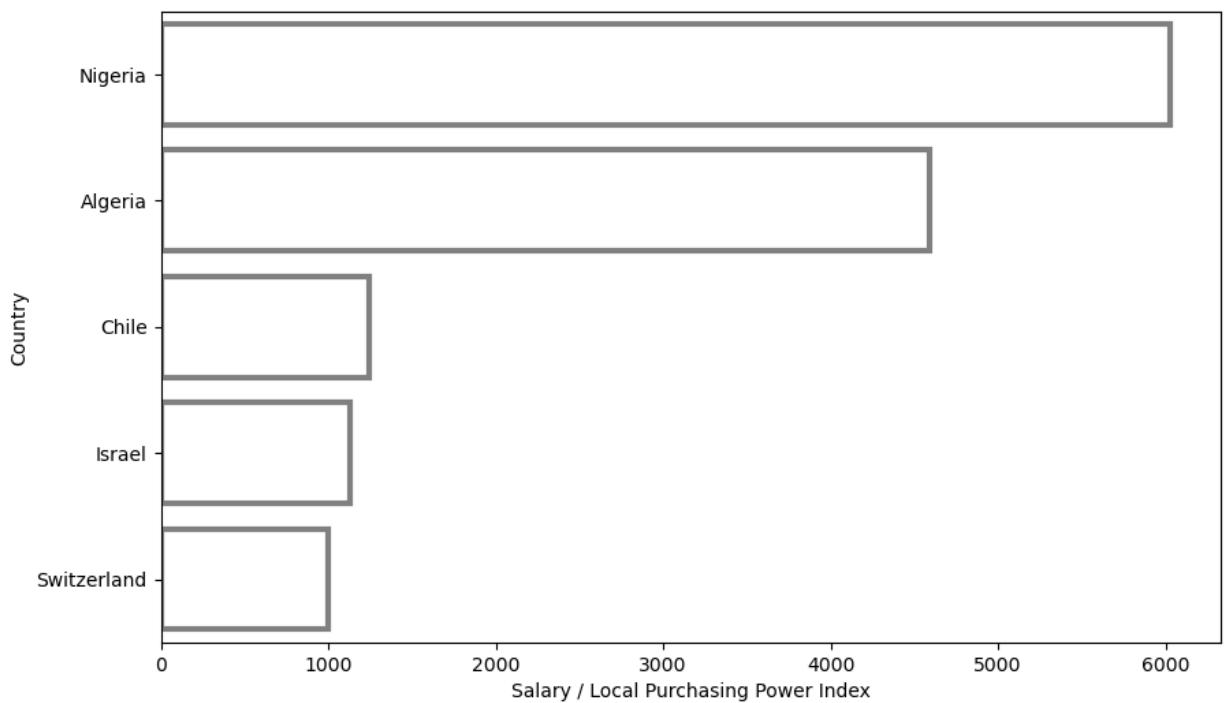
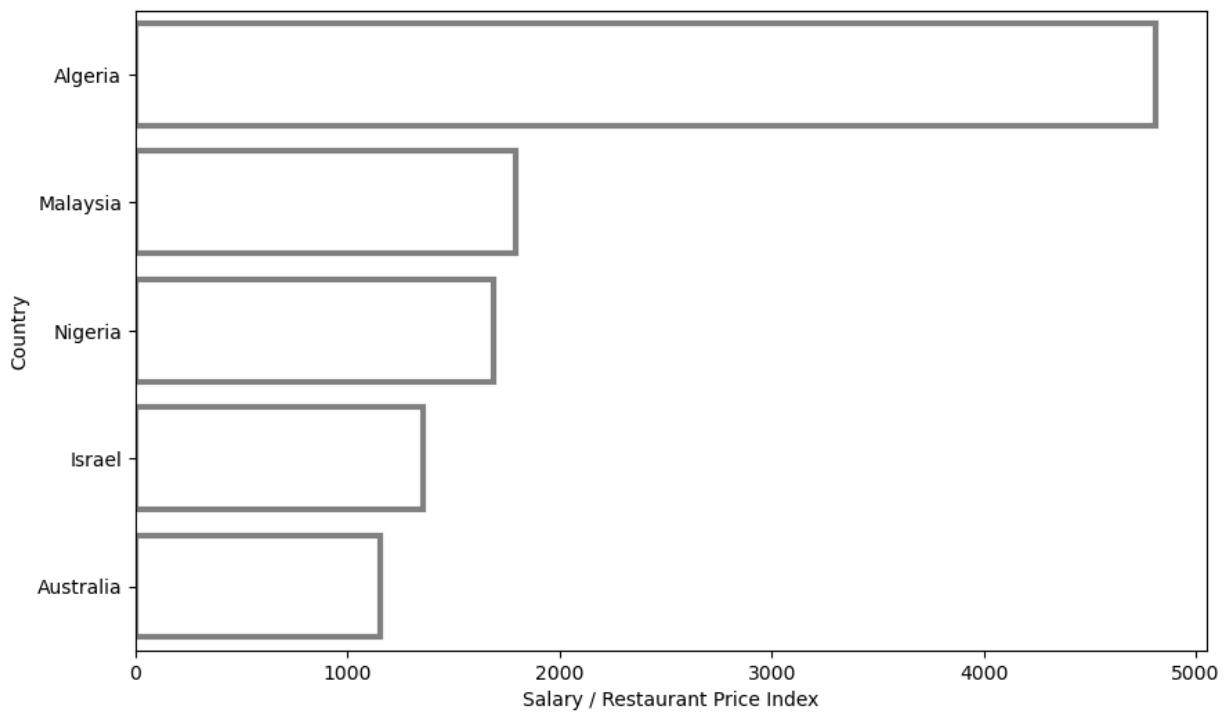
```
top_countries_by_index = {}

for column in index_columns:
    ratio_column = f'Salary / {column}'
    grouped_data[ratio_column] = grouped_data['Salary (USD)'] / grouped_data[column]
    top_countries_by_index[column] = grouped_data.nlargest(5, ratio_column)

for column, data in top_countries_by_index.items():
    create_bar_plot(data, f'Salary / {column}', 'Country', color_palette)
```







```
In [ ]: print("Algeria seems to be the place to be. There are remote jobs there ")
```

```
In [60]: print(ds_salaries[ds_salaries['company_location'] == 'DZ']) #algeria
print(ds_salaries[ds_salaries['company_location'] == 'CH']) #switzerland
print(ds_salaries[ds_salaries['company_location'] == 'NG']) #nigeria
```

#here are just some examples of jobs from 3 countries.

```

        Unnamed: 0  work_year experience_level employment_type      job_title  \
487          487           2022                  EN                  PT Data Scientist

        salary salary_currency  salary_in_usd employee_residence  remote_ratio  \
487  100000          USD            100000                  DZ                 50

        company_location company_size
487             DZ              M
        Unnamed: 0  work_year experience_level employment_type  \
213          213           2021                  EN                  FT
518          518           2022                  MI                  FT

        job_title  salary salary_currency  salary_in_usd  \
213 Big Data Engineer   435000            INR         5882
518 Data Scientist     115000            CHF        122346

        employee_residence  remote_ratio company_location company_size
213             IN              0                  CH              L
518             CH              0                  CH              L
        Unnamed: 0  work_year experience_level employment_type      job_title  \
38            38           2020                  EN                  FT Data Analyst
116          116           2021                  MI                  FT Data Scientist

        salary salary_currency  salary_in_usd employee_residence  remote_ratio  \
38    10000          USD            10000                  NG                100
116   50000          USD            50000                  NG                100

        company_location company_size
38             NG              S
116             NG              L

```

Conclusion:

As you can see Data Scientist and Data related field's have many jobs available throughout the world. It is possible to work remote and make large salary from a company in Switzerland and live in Isreal. There are many combinations. I could not provide a definitive suggestion as my analysis lack many things.

Things I could have improved on:

- I could have included the united states two letter states entries but I filtered them out and did not include them in the average for United Stated merged data.
- I could have grouped the entries rather than writing code to mean as I could have seen how many entries made the average as I could also plot how many jobs were from a specific region
- I could have also doing a z score test to give an overall score to see which country's salary would go the further but I chose to show all the different countries for all the different indexes.

I apologize for the delayed submission and the lack of analysis for this project. I tried to highlight everything I have learned in the course but due to lack of time I could not dive in further. I would appreciate any feedback on

how to improve for my future projects and when I get the chance I will take your SQL course.